Early Disease Detection

Predicting Diabetes Risk Using Data Science for Early Disease Detection

Prepared by: Dorka Swieton
Date: January 23, 2026
Program: Tomorrow University / IC Data Science Fundamentals Capstone

Introduction:
Diabetes affects more than 537 million adults worldwide, and nearly one in two cases remains undiagnosed, often leading to severe complications and global healthcare costs exceeding USD 960 billion annually. The core problem is that traditional screening methods frequently detect diabetes too late, limiting the effectiveness of preventive care. This proposal recommends implementing a data science–based risk prediction system that analyzes patient data such as age, body mass index, blood glucose levels, and family history to identify high-risk individuals earlier. By enabling proactive intervention, healthcare providers can improve patient outcomes, reduce treatment costs, and use resources more efficiently. It is therefore recommended that healthcare organizations adopt data-driven prediction tools to support early detection and strengthen preventive healthcare strategies.

Based on the analysis of the Diabetes Risk Factors dataset, a key proposed solution is the development and implementation of a data science–driven diabetes risk prediction system to support early identification and preventive healthcare interventions. The dataset highlights several significant predictors of diabetes, including age, body mass index (BMI), high blood pressure, high cholesterol levels, and lifestyle-related factors such as physical inactivity. Notably, the analysis indicates that individuals between the ages of 40 and 60 represent a particularly high-risk group, underscoring the importance of targeted risk assessment within this population segment.

To address this challenge, healthcare organizations can implement a machine learning–based risk assessment model that analyzes patient health data and generates individualized risk scores. This system would enable healthcare providers to proactively identify individuals with an elevated probability of developing diabetes, even before clinical symptoms become apparent. Based on these insights, providers can recommend appropriate preventive measures, including lifestyle modifications, regular monitoring, and early medical intervention.

In addition, the integration of such a predictive system into electronic health records would support data-driven clinical decision-making, allowing healthcare professionals to allocate preventive resources more efficiently and prioritize high-risk patients. This proactive approach has the potential to improve patient outcomes, reduce the incidence of advanced diabetes complications, and significantly lower long-term healthcare costs.

1. Business Case – The Problem & Opportunity

The Problem

Diabetes represents one of the most significant and rapidly growing public health challenges worldwide. According to global health data, hundreds of millions of adults are currently living with diabetes, and a substantial proportion of cases remain undiagnosed. A key issue is that diabetes often develops gradually and without obvious symptoms in its early stages, resulting in delayed diagnosis and missed opportunities for prevention.

Late detection has serious consequences. Patients diagnosed at advanced stages face a significantly higher risk of complications, including cardiovascular disease, kidney failure, nerve damage, and vision impairment. These complications not only reduce quality of life but also require complex, long-term, and costly medical treatment. For healthcare providers and systems, this leads to increased financial burden, inefficient use of medical resources, and reduced capacity to focus on preventive care.

Traditional screening approaches rely primarily on periodic medical examinations and general clinical guidelines. However, these methods may not fully capture individual risk profiles or identify high-risk patients early enough. As a result, healthcare providers often operate reactively, treating the disease after it has already progressed rather than preventing it.

The Opportunity

Advances in data science and machine learning present a significant opportunity to transform diabetes prevention through early risk identification. Healthcare systems already collect large volumes of patient data, including demographic information, medical history, and clinical measurements. This data can be leveraged to identify patterns and risk factors associated with the development of diabetes.

By implementing a data science–based risk prediction system, healthcare providers can analyze patient data to identify individuals who are at elevated risk before the disease fully develops. This enables earlier intervention through preventive measures such as lifestyle recommendations, closer monitoring, and early treatment.

This approach creates value in several ways. It improves patient outcomes by reducing the likelihood of severe complications, lowers long-term healthcare costs by preventing advanced disease, and enhances operational efficiency by enabling more targeted and proactive care.

Ultimately, the integration of data science into diabetes risk assessment represents a strategic opportunity for healthcare organizations to shift from reactive disease management to proactive prevention, improving both clinical effectiveness and economic sustainability.

2. Data Source Evaluation

Source: Kaggle – Diabetes Risk Factors Dataset

The primary data source used in this project is the Diabetes Risk Factors dataset available on Kaggle, which is based on the Behavioral Risk Factor Surveillance System (BRFSS), a large-scale health survey conducted by the U.S. Centers for Disease Control and Prevention (CDC). This dataset contains over 250,000 individual survey responses and includes approximately 21 health-related variables, making it highly suitable for analyzing diabetes risk factors and building predictive models.

Relevance and Strengths of the Dataset

This dataset is highly relevant to the business problem of early diabetes detection because it includes a wide range of variables that are strongly associated with diabetes risk. These include:

- Age
- Body Mass Index (BMI)
- High blood pressure
- High cholesterol
- Physical activity
- Smoking status
- General health condition

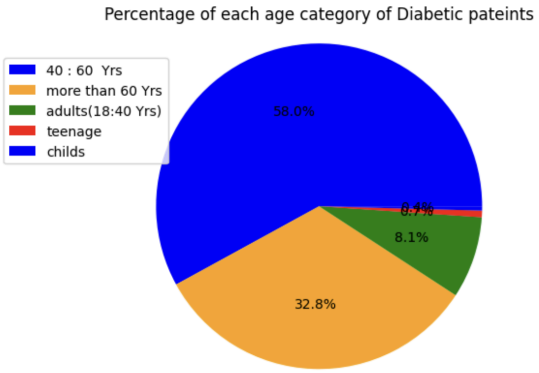

These indicators are widely recognized as important predictors of diabetes and are commonly used in machine learning models for early risk detection.

One of the key strengths of this dataset is its large sample size, which improves the reliability and statistical validity of the analysis. Larger datasets allow machine learning models to identify patterns more accurately and reduce the likelihood of random error.

Another important strength is that the dataset is based on real-world healthcare survey data collected by a reputable public health organization. The BRFSS survey is the world's largest continuously conducted health survey, designed to monitor behavioral risk factors related to chronic diseases such as diabetes.

This increases the credibility and reliability of the data.

Data Quality and Limitations

Despite its strengths, the dataset has several limitations that must be considered.

Self-Reported Data

One major limitation is that much of the information in the dataset is self-reported by survey participants. This introduces the risk of:

- Inaccurate reporting
- Recall bias
- Underreporting of unhealthy behaviors

This may affect the accuracy of predictions.

Geographic Limitation

The dataset is based on survey responses from the United States population. As a result, the findings may not fully represent populations in other countries due to differences in:

- lifestyle
- healthcare systems
- genetic factors

This limits generalizability.

Missing or Imbalanced Data

Like many real-world healthcare datasets, some variables may be imbalanced, meaning that there are more non-diabetic cases than diabetic cases. This can bias predictive models if not properly handled.

However, this limitation can be addressed using standard data science techniques such as:

- data balancing
- sampling
- preprocessing

Data Gaps

While the dataset includes many important health indicators, it does not include certain clinical measures such as:

- laboratory blood test results
- genetic data
- detailed medical history

Including these variables could improve prediction accuracy.

3. Ethical Risk Assessment & Regulatory Alignment

The implementation of a data science–based diabetes risk prediction system presents significant opportunities to improve preventive healthcare. However, it also introduces important ethical risks that must be carefully addressed to ensure responsible, fair, and trustworthy use of data and artificial intelligence.

3.1 Ethical Risk 1: Data Privacy and Confidentiality

One of the primary ethical concerns is the protection of patient privacy. Diabetes risk prediction models rely on sensitive personal and health-related information, such as age, medical history, and lifestyle factors. Unauthorized access, data breaches, or misuse of this information could result in harm to individuals, including discrimination, loss of privacy, or reputational damage.

Mitigation Strategy

To address this risk, healthcare organizations must implement strong data protection measures, including:

- anonymization or pseudonymization of personal data
- secure data storage and encryption
- restricted access to authorized personnel only

These safeguards help ensure that patient information remains confidential and protected.

3.2 Ethical Risk 2: Bias and Fairness

Another important ethical risk is algorithmic bias. If the dataset used to train the prediction model does not fully represent all population groups, the model may produce less accurate predictions for certain individuals. This could lead to unequal treatment, misclassification, or unfair healthcare decisions.

For example, if certain demographic groups are underrepresented in the dataset, the model's predictions may not be equally reliable for those populations.

Mitigation Strategy

This risk can be mitigated by:

- using diverse and representative datasets
- regularly evaluating model performance across different population groups
- continuously updating and improving the model

Ensuring fairness and accuracy across all groups is essential for ethical implementation.

3.3 Ethical Risk 3: Over-Reliance on Automated Decision-Making

There is also a risk that healthcare professionals may rely too heavily on the model's predictions without applying their own clinical judgment. This could result in inappropriate decisions if the model's predictions are incorrect or incomplete.

Mitigation Strategy

The prediction system should be used as a decision-support tool, not a replacement for medical professionals. Final decisions should always remain under human supervision.

This ensures that data science enhances, rather than replaces, professional expertise.

Regulatory Alignment

The proposed solution must comply with relevant data protection and artificial intelligence regulations, particularly within the European Union.

General Data Protection Regulation (GDPR)

The GDPR is the primary legal framework governing the use of personal data in the European Union. It requires that personal data be:

- processed lawfully, fairly, and transparently
- collected for specific and legitimate purposes
- protected against unauthorized access

Under GDPR, healthcare organizations must also ensure that patients are informed about how their data is used and that their privacy rights are respected.

EU Artificial Intelligence Act

The EU AI Act establishes rules for the responsible use of artificial intelligence, particularly in high-risk areas such as healthcare.

Under the EU AI Act, AI systems used in healthcare must ensure:

- transparency
- accuracy
- safety
- human oversight

The proposed diabetes risk prediction system would be classified as a high-risk AI system, meaning it must meet strict requirements for risk management, monitoring, and accountability.

4. Structural Elements and Data Quality Issues of the Diabetes Risk Factors Dataset

The Diabetes Risk Factors dataset is a structured, tabular dataset containing over 250,000 records, where each row represents an individual and each column represents a health or demographic variable. The dataset includes a clear target variable (diabetes diagnosis) and multiple predictor variables, such as age, BMI, blood pressure, cholesterol, and lifestyle factors. This structure makes it suitable for supervised machine learning to predict diabetes risk. However, several structural issues affect its interpretation and use.

First, the dataset has class imbalance, with significantly more non-diabetic than diabetic cases. This can cause machine learning models to favor predicting the majority class,

reducing their ability to correctly identify high-risk individuals. This directly impacts the business goal of early detection, as some at-risk patients may be missed.

Second, many variables are based on self-reported data, which may contain inaccuracies due to recall errors or misreporting. This reduces data reliability and may lower prediction accuracy, affecting the trustworthiness of the model's recommendations.

Third, the dataset lacks detailed clinical measurements, such as laboratory blood glucose values. As a result, the model relies on indirect indicators, making it more suitable for risk screening rather than medical diagnosis.

These limitations highlight the importance of careful data preprocessing, validation, and human oversight to ensure the data science solution provides accurate and responsible support for early diabetes detection.

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level	diabetes
0	Female	80.0	0	1	never	25.19	6.6	140	0
1	Female	54.0	0	0	No Info	27.32	6.6	80	0
2	Male	28.0	0	0	never	27.32	5.7	158	0
3	Female	36.0	0	0	current	23.45	5.0	155	0
4	Male	76.0	1	1	current	20.14	4.8	155	0

5. Business Value Proposition

The implementation of a data science–based diabetes risk prediction system offers significant value for healthcare organizations by enabling earlier identification of high-risk individuals and supporting preventive care strategies. By analyzing patient data such as age, BMI, hypertension status, HbA1c level, blood glucose level, and smoking history, the system can generate accurate risk assessments that help healthcare providers intervene before the disease progresses.

One of the primary business benefits is cost reduction. Treating advanced diabetes and its complications is substantially more expensive than preventing the disease. By identifying at-risk patients early, healthcare providers can recommend lifestyle changes and early treatment, reducing the need for costly long-term care, hospitalizations, and complication management.

In addition, the system improves operational efficiency by helping healthcare professionals prioritize patients based on their risk level. This allows for more effective allocation of medical resources, ensuring that preventive efforts are focused where they will have the greatest impact.

The solution also enhances clinical decision-making by providing data-driven insights to support healthcare professionals. Rather than relying solely on periodic testing or general guidelines, providers can use predictive analytics to make more informed and proactive decisions.

Finally, the system contributes to better patient outcomes and improved quality of care. Early detection and prevention reduce the likelihood of severe complications, improve patient quality of life, and strengthen trust in healthcare services.

Overall, investing in a diabetes risk prediction system creates both economic value and healthcare value, supporting more efficient, proactive, and sustainable healthcare delivery.

6. Risk Mitigation and Operations

The successful implementation of a data science–based diabetes risk prediction system requires careful management of operational and ethical risks to ensure accurate, safe, and reliable use in healthcare settings.

One key operational risk is data quality and accuracy. If the input data is incomplete, outdated, or inaccurate, the model's predictions may be unreliable. To mitigate this risk, healthcare organizations should implement data validation procedures, regularly update the dataset, and monitor the model's performance over time. Continuous evaluation ensures that the system maintains accuracy and remains relevant to current patient populations.

Another important risk is over-reliance on the prediction system. There is a possibility that healthcare professionals may depend too heavily on automated predictions without applying their clinical expertise. To address this, the system should be implemented as a decision-support tool rather than a replacement for medical professionals. Final clinical decisions must remain under human supervision to ensure responsible use.

Data privacy and security risks must also be carefully managed, as the system uses sensitive personal and health information. Healthcare organizations should implement strong security measures, including data encryption, secure storage, and restricted access to authorized personnel only. These measures ensure compliance with data protection regulations and maintain patient trust.

From an operational perspective, successful deployment requires integration into existing healthcare workflows, such as electronic health record systems. Healthcare staff should receive appropriate training to understand how to interpret and use the model's predictions effectively.

Finally, ongoing maintenance and monitoring are essential. The system should be regularly reviewed, tested, and updated to ensure continued accuracy, fairness, and compliance with regulatory standards.

By implementing these risk mitigation and operational strategies, healthcare organizations can ensure that the diabetes risk prediction system operates safely, effectively, and responsibly while delivering its intended benefits.

7. Strategic Coherence and Final Recommendation

Problem

Diabetes is a widespread chronic disease that is often diagnosed too late, after significant health complications have already developed. Traditional screening methods rely on periodic testing and general risk factors, which may not identify high-risk individuals early enough. This reactive approach leads to poorer patient outcomes, higher treatment costs, and inefficient use of healthcare resources. Therefore, there is a clear need for a proactive, data-driven solution that enables earlier identification of individuals at risk of developing diabetes.

Data Strategy

To address this problem, this project uses the Diabetes Prediction Dataset from Kaggle, which contains structured patient-level data, including age, BMI, hypertension status, heart disease, smoking history, HbA1c level, and blood glucose level. These variables are well-established predictors of diabetes risk. By applying supervised machine learning techniques to this dataset, patterns and relationships between health indicators and diabetes outcomes can be identified. Proper data preprocessing, validation, and ongoing monitoring ensure that the model produces accurate and reliable predictions to support early risk detection.

Ethics and Regulation

The use of health data requires strict adherence to ethical standards and regulatory frameworks. Key ethical considerations include protecting patient privacy, ensuring data security, and preventing algorithmic bias. This solution aligns with the General Data Protection Regulation (GDPR) by ensuring that personal data is handled securely and used only for legitimate healthcare purposes. In addition, the system complies with the EU AI Act, which requires transparency, human oversight, and risk management for AI systems used in healthcare. The model is designed to support, not replace, healthcare professionals, ensuring that final decisions remain under human control.

Solution

The proposed solution is the implementation of a data science–based diabetes risk prediction system that analyzes patient health indicators and generates individual risk assessments. Using machine learning, the system identifies patients with a higher probability of developing diabetes and provides healthcare providers with actionable insights. This enables earlier intervention, including preventive care, lifestyle recommendations, and closer monitoring. The system can be integrated into existing healthcare platforms to support routine clinical workflows and improve preventive care strategies.

Business Value

This solution provides significant value for healthcare organizations by improving efficiency, reducing costs, and enhancing patient outcomes. Early detection helps prevent severe complications, which lowers long-term treatment expenses and reduces the burden on healthcare systems. The system also enables more efficient allocation of medical resources by focusing preventive efforts on high-risk individuals. Additionally, data-driven decision-making improves the quality of care and strengthens trust in healthcare services. Overall, the solution supports more proactive, efficient, and sustainable healthcare delivery.

Final Recommendation

Based on the analysis, it is strongly recommended that healthcare organizations invest in the development and implementation of a diabetes risk prediction system using data science. The available data provides a reliable foundation for predictive modeling, and appropriate ethical and regulatory safeguards can ensure responsible use. This solution aligns with both clinical and business objectives by improving early detection, reducing healthcare costs, and enhancing patient care. Implementing this system represents a strategic opportunity to shift from reactive treatment to proactive prevention, creating long-term value for both healthcare providers and patients.